

# Settlement Exception Management Policy

|                   |                              |                    |                       |
|-------------------|------------------------------|--------------------|-----------------------|
| Policy ID         | POL-SET-001                  | Version            | 1.0                   |
| Status            | SUPERSEDED                   | Effective Date     | 2024-01-01            |
| Review / End Date | 2024-12-31                   | Supersedes         | None                  |
| Policy Owner      | Global Settlement Operations | Business Unit      | SETTLEMENTS           |
| Jurisdiction      | GLOBAL                       | Classification     | INTERNAL CONFIDENTIAL |
| Access Scope      | ROLE_BASED                   | Synthetic Document | YES                   |

**Document control:** The applicable policy is the version whose effective-date range contains the business date and whose status is ACTIVE. A higher future version must not be used for current decisions.

## 1. Policy Objective

This policy supports production-grade asset-servicing controls and AI-assisted operations. Decisions must be grounded in authorized sources, linked to the correct business date, and preserved with sufficient evidence for audit and independent review.

## 2. Purpose and Scope

Defines controls for identifying, prioritizing, investigating, escalating, and resolving post-trade settlement exceptions across supported markets. Every material action must include user identity, source-system reference, policy citation, timestamp, and escalation evidence when applicable.

## 3. Exception Classification

Classify by client impact, market cutoff risk, monetary exposure, regulatory sensitivity, and availability of a safe workaround. Every material action must include user identity, source-system reference, policy citation, timestamp, and escalation evidence when applicable.

## 4. Investigation Controls

Validate trade economics, account status, standing settlement instructions, counterparty data, available cash or securities, and market status. Every material action must include user identity, source-system reference, policy citation, timestamp, and escalation evidence when applicable.

## 5. Escalation and Closure

Escalate based on priority and cutoff risk. Closure requires reason code, evidence, ownership, and checker approval where required. Every material action must include user identity, source-system reference, policy citation, timestamp, and escalation evidence when applicable.

## 6. Current Version-Specific Rules

| Control                  | Requirement in v1.0      |
|--------------------------|--------------------------|
| P1 acknowledgement       | 120 minutes              |
| P2 acknowledgement       | 240 minutes              |
| Market-cutoff escalation | 60 minutes before cutoff |
| Evidence retention       | 5 years                  |

**Control principle:** live operational status comes from governed source-system tools; policy documents provide rules, required actions, and escalation criteria.


## Required Evidence and Data Controls

| Area            | Minimum production requirement                                                                |
|-----------------|-----------------------------------------------------------------------------------------------|
| Identity        | Authenticated user, role, business unit, and permitted account/client scope                   |
| Source data     | Current timestamp, source-system identifier, validated business keys, and freshness status    |
| Policy evidence | Policy ID, applicable version, effective range, section, and citation location                |
| AI processing   | PII masking result, prompt version, model/deployment, retrieval set, and confidence signals   |
| Audit           | Trace ID, tool-call records, response outcome, escalation decision, and reviewer where needed |

## Key Definitions

**Applicable version:** the authorized version active for the business date. **Governed tool:** an approved interface that validates identity, permissions, parameters, and audit evidence before reading an operational system. **Material case:** an item whose amount, age, deadline, client impact, repeated occurrence, or regulatory sensitivity requires enhanced review.

Operational Control Flow



Roles and Responsibilities

| Role                     | Primary responsibility                                                                                |
|--------------------------|-------------------------------------------------------------------------------------------------------|
| Operations Analyst       | Investigate assigned items, validate evidence, document actions, and escalate when required.          |
| Team Lead / Checker      | Review sensitive or high-value actions, validate policy use, and approve controlled overrides.        |
| Policy Control           | Maintain versions, effective dates, interpretation guidance, and formal exceptions.                   |
| AI / Technology Platform | Maintain access controls, approved integrations, retrieval quality, monitoring, and failure handling. |
| Information Security     | Define masking, secrets, privileged access, incident response, and monitoring requirements.           |

Audit, Privacy, and AI Requirements

- Record policy ID, version, page or section, query timestamp, user identity, access scope, and decision outcome.
- For live lookups, record MCP tool name, approved parameters, authorization result, response timestamp, and source-system reference.
- Mask PII and confidential fields in prompts and traces unless the value is explicitly required and permitted for the task.
- Scope semantic-cache entries by user authorization, business unit, policy version, and effective date; invalidate them when policy status changes.
- Route conflicting evidence, unavailable tools, restricted access, and low-confidence or high-risk cases to human review.

Test Scenarios

| Scenario                                  | Expected behavior                                                   |
|-------------------------------------------|---------------------------------------------------------------------|
| Older policy chunk ranks higher           | Use the applicable ACTIVE version and suppress superseded guidance. |
| Future policy has highest version         | Do not use it before its effective date.                            |
| User lacks account permission             | Return an authorization-safe response without restricted details.   |
| Cached answer uses older version          | Bypass or invalidate cache and regenerate with current policy.      |
| Live status differs from document example | Use governed live data for status and policy only for rules.        |

Monitoring and Acceptance Criteria

| Control metric          | Expected production behavior                                                                        |
|-------------------------|-----------------------------------------------------------------------------------------------------|
| Policy-version accuracy | Current applicable version selected for the business date; future and superseded versions excluded  |
| Retrieval quality       | Relevant section appears in top results and citations match returned evidence                       |
| Tool governance         | Every live lookup is authorized, parameter-validated, rate-limited, and auditable                   |
| Privacy                 | Sensitive values are minimized and redacted from cache, prompts, traces, and logs                   |
| Response quality        | Answer is grounded, cites evidence, distinguishes live facts from policy, and escalates uncertainty |

Implementation Notes for the AI Copilot

The Retrieval Agent searches approved policy content using hybrid retrieval and metadata filters. The Policy Validation Agent confirms status, effective dates, and conflicts. Operational agents use MCP tools for current records. The Orchestrator aggregates all agent outputs and sends one structured evidence package to the Response Module for grounded answer generation, citations, masking, confidence assessment, and escalation.

End of controlled synthetic document. Structured for OCR/layout extraction, document-aware chunking, metadata enrichment, version-aware retrieval, citations, PII controls, semantic-cache invalidation, and policy validation.