

AI Copilot Human Escalation and Audit Review

SOP ID	SOP-AI-001	Version	1.0
Status	ACTIVE	Effective Date	2026-01-01
Parent Policy	POL-SEC-001	Business Unit	ENTERPRISE
Classification	INTERNAL CONFIDENTIAL	Access Scope	ROLE_BASED
Owner	Asset Servicing Operations	Review Date	2026-12-31
Synthetic Document	YES	Maker-Checker	WHEN REQUIRED

Purpose

This SOP defines a repeatable, auditable workflow for authorized asset-servicing operations. It separates policy guidance from current operational status and requires governed access to live data.

End-to-End Procedure



1. Authenticate and classify request

Confirm identifiers, business date, user role, and access scope before processing.

2. Mask sensitive data and check cache scope

Use only approved interfaces and minimum necessary data. Record source references and timestamps.

3. Collect agent and MCP evidence

Compare live evidence with the applicable policy version; never substitute a future or superseded version.

4. Assess confidence, conflicts, and risk

Apply priority, materiality, cutoff, and maker-checker rules relevant to the case.

5. Generate grounded response or human escalation

Record outcome, reason code, citations, confidence, approval, and any escalation or follow-up.

Control Checkpoints

Checkpoint	Required evidence
Authorization	Authenticated identity, role, business unit, account or client scope
Policy	Policy ID, applicable version, effective date, section or rule
Operational data	MCP tool, validated parameters, source-system timestamp
Decision	Reason, action, confidence, citations, checker or escalation
Privacy	Masking status and trace/log redaction outcome

Inputs, Outputs, and Ownership

Item	Definition
Inputs	Authorized user request, validated business identifiers, current operational records, applicable policy evidence
Outputs	Finding, evidence summary, required action, owner, deadline, citations, confidence, escalation status
Primary owner	Assigned operations analyst or specialist queue
Checker / approver	Team lead, policy control, information security, or designated risk owner based on the case
System record	Case-management or audit record linked to source-system and AI trace identifiers

Exception and Escalation Handling

- Authorization denied: do not retry with broader privileges; return a controlled access response and log the denial.
- Applicable policy cannot be established: stop automated recommendation and route to Policy Control.
- Live data and cache disagree: bypass the cache and refresh through the approved MCP tool.
- PII appears in free text: mask before LLM and tracing unless explicitly approved for the task.
- Tool failure, conflicting evidence, high financial exposure, regulatory sensitivity, or low confidence: create a human-review package.

Human Review Package

Field	Content
Request summary	What the user asked and the operational identifiers involved
Evidence summary	Retrieved policy passages and governed live-system results
Risk and conflict	Missing data, conflicting versions, authorization limits, SLA or materiality risk
Recommended next step	Safe action, required owner, and deadline
Audit references	Session, trace, tool-call, source-system, and citation identifiers

Quality and Completion Checklist

- Correct identifiers and current source data were used.
- The applicable ACTIVE policy version was verified.
- User authorization and ACL scope were enforced.
- PII was minimized and masked in prompts, cache, traces, and logs.
- Citations and tool references allow independent reproduction.
- The final response is grounded, concise, and escalated when required.

Example Outcome Format

Finding: concise operational result. **Evidence:** live source-system references plus current policy citations. **Required action:** owner and deadline. **Risk:** SLA, client, market, regulatory, or access concern. **Escalation:** required or not required.

Audit Record Fields

Field group	Examples
Request	Session ID, user ID, role, business unit, normalized intent, identifiers
Retrieval	Search query, filters, document IDs, policy versions, page/section citations, reranking scores
MCP tools	Tool name, validated parameters, permission result, source timestamp, response status
Decision	Agent findings, conflicts, confidence, guardrail results, escalation reason
Response	Prompt version, final citations, masking status, latency, token use, feedback status

Operational Quality Targets

The workflow should return current and authorized facts, select the correct effective policy, preserve source references, avoid exposing restricted data, and provide enough evidence for a human reviewer to reproduce the result. Failures must be visible through structured logs, traces, metrics, and alerting rather than silently hidden.

End of controlled synthetic SOP. Designed for Document Intelligence, section-aware chunking, secure retrieval, MCP tool use, agent orchestration, response generation, and audit testing.